

MPC Controller

Andreas Gotterba

1. Edited Files

agotterba_mpc_writeup.pdf

This Document

src/

main.cpp

main program

MPC.h

MPC controller header file

MPC.cpp

MPC controller implementation

Videos

50 mph lap:

<https://youtu.be/q0zoL7J5ZWo>

2. Model Description

2.1 State

The state variables are:

- x: The x coordinate of the car
- y: The y coordinate of the car
- psi: The angle the car is pointing
- v: The speed of the car
- cte: The cross track error of the car, relative to the computed (polynomial) centerline
- epsi: The error in the angle of the car, relative to the computed angle of the centerline

Note that all measurements are made relative to the car's current position. So in the car's current state, x, y, and psi are all 0. These state variables take on different values for future time-steps, as the model predicts what each of these values will be *relative to the car's current position*.

2.2 Actuators

The actuators are:

- sangle: The steering angle of the car (processed in radians, and then converted to '25 degrees = 1.0' at the end
- acc: The acceleration (throttle) of the car. The model is ideal (so it really considers this acceleration), but there is resistance/friction/drag in the simulator. Thus, when the model targets 50mph, it ends up running at around 48.5 mph.

2.3 Update Equations

For each time-step, the model predicts the values of the state variables, based on the previous state and the values of the actuators that were applied. The fundamental equations are:

$$\begin{aligned}x[t+1] &= x[t] + v[t] * \cos(\text{psi}[t]) * dt \\y[t+1] &= y[t] + v[t] * \sin(\text{psi}[t]) * dt \\\text{psi}[t+1] &= \text{psi}[t] - (v[t]/L_f) * \text{sangle}[t] * dt \\v[t+1] &= v[t] + \text{acc}[t] * dt \\\text{cte}[t+1] &= f(x[t]) - y[t] + v[t] * \sin(\text{epsi}[t]) * dt \\\text{epsi}[t+1] &= \text{psi}[t] - \text{psides}(x[t]) + (v_0/L_f) * \text{sangle}[t] * dt\end{aligned}$$

Most of the operands are the state and actuators described above; the additional parameters are:

Lf: The length from the front axle of the car to its center of gravity. Supplied as part of the starter code, from the result of a controlled simulation.

$f(x[t])$: this function returns the y coordinate for the centerline of the road at a given x coordinate (and I use the x coordinate for time t). This could be any means of figuring out the y coordinate; in this implementation, I'm using the polynomial model from the waypoints received from the simulator, so I simply evaluate the polynomial at $x[t]$.

$psides(x[t])$: this function returns the angle of the road's centerline at a given x coordinate (and I pass in the x coordinate for time t). This could be any means of finding the angle; in this implementation, I'm using the polynomial model from the waypoints received from the simulator; I compute the derivative of the polynomial, then take the arctangent of that to find the angle.

The actual equations I'm using are in lines 180 through 189 of MPC.cpp. Note that the code is a bit more involved than what I've shown here. Instead of the raw update equations, these are implemented as constraints. On its own, the model would think it could pick any value for $x[t+1]$ that it likes. But I've already constrained $fg[1+param_start+t]$ to be zero. Therefore, these equations tell it that $x[t+1]$ minus (the equation I wrote above) should be zero, thus constraining the value of $x[t+1]$ that it picks to be equal to the update equation written above.

3. Timestep Length and Elapsed Duration

I selected my values for elapsed duration (or rather, number of timesteps N) and the timestep length (dt) iteratively, but with a few guiding ideas:

1. I watched the distance that was being predicted ahead by plotting all the predicted values using the `mpc_x` and `mpc_y` keys in the json message, and compared it to the range of the waypoints supplied by the simulator. Predicting beyond the waypoints isn't useful, since the polynomial it's trying to fit is no longer matching the road. This distance varies around the track (becoming quite short in the tight turns), but as N and dt are constant, I chose a distance I thought struck a good compromise around the track.

2. I actually set T (the amount of time to look ahead) to a fixed time and compute N from T and dt . That way, if I decide to change dt but want the modeled amount of time to stay the same, N automatically adjusts. I found that looking half a second ahead wasn't enough, but a full second is fine.
3. I want dt to be small enough to maintain accuracy, but not so small as it has a noticeable impact on performance. When I set T to 1 second and dt to 20ms (N was therefore 50), I did see worse performance. Things are much better with $T = 1$ second and $dt = 50$ ms ($N = 20$).
4. As I'll discuss later, I'd like the latency to be a multiple of dt , since that simplifies how I'm handling latency (means I don't need any interpolation).

4. Polynomial Fitting and MPC Preprocessing

When I started the project, I tried keeping my calculations in absolute coordinates, imagining that if I was getting the waypoints from localization, that would be the frame of reference I was last in. I soon realized this was problematic: since the path is being approximated as a polynomial, what would happen when the car is headed in the y direction, and the correct line is no longer a function. The path would go crazy, as it tries to fit in the wrong direction. Once I switched to working in the car's coordinates, everything became easier. I read the waypoints into `raw_ptsx` and `raw_pty`, then transform them to the car's coordinates (using the reported x, y , and ψ values). Then I know that the car is at 0,0 and ψ is 0 (in the car's reference frame). I can call the supplied `polyfit` function on the transformed waypoints, and know that `cte` is the constant coefficient and `epsi` is $-1 * \text{the arctangent of the linear coefficient}$.

From there, I know the current state is $(x, y, \psi, v, cte, epsi) = (0, 0, 0, v, cte, epsi)$ where v is reported from the simulator while `cte` and `epsi` are calculated from the polynomial as I described, and I can call `Solve` with that state and the polynomial's coefficients as inputs.

Once the `steer_value` and `throttle_value` are returned, I scale the `steer_value` (which was calculated to be between ± 0.436 radians) to range between ± 1 . The throttle value was already in its ± 1 range, so it doesn't need to be scaled. The steering and throttle values are then ready to be sent to the simulator.

5. Latency

There was a lot of talk on the forums about predicting where the car will be in 100ms, by just assuming it continues with the steering angle and throttle that it currently has. That seemed to me like a poor man's model predictor- why should I use that when I can remember what values of `acc` and `sangle` were

sent in previous timesteps, and have already implemented a heavy-duty model predictor that gets run every iteration?

So instead, I compute `n_latency`: the number of timesteps that are consumed by the latency; in this case `n_latency` is 2. I've added `n_latency` as a variable of the MPC object, as well as `sangle_hist` and `acc_hist`, which store the previous values of steering angle and throttle; they get initialized to zeros in the `init` function.

When the model makes its prediction, it constrains the values of `sangle` and `throttle` for the first two timesteps to be those that have already been sent (at lines 289 and 306 of `MPC.cpp`). This is similar to how it normally constrains these to ± 0.436 and ± 1 ; just for these time steps, they are constrained to be exactly the values that were sent.

Then when reporting the values to use next, instead of sending the values from the first timestep of the model (`sangle[0]` and `acc[0]`), I send `sangle[n_latency]` and `acc[n_latency]` (in this case `sangle[2]` and `acc[2]`), and push those values into the `sangle_hist` and `acc_hist` queues, to be ready for the next iteration.

I found this system to work extremely well- the car is very stable, and stays centered in the lane. I confirmed that when I enabled the latency but did not correct for it, it would oscillate across the road.

I tried adding even more latency to the computation (the `effective_latency`) while the `applied_latency` remained 100ms, to account for the additional time to compute the new values, and get a new request from the simulator. But I found that while I could measure this loop delay with `system_clock` commands, I couldn't tell when the new values were applied within the simulator. I could probably nail it down by trial and error, but found that performance was fine when I simply set `effective_latency` equal to the `applied_latency`.

Note that because I use the latency to compute `n_latency`, I set it as a variable (`applied_latency`) at line 15 of `main.cpp`, instead of hard coding it at the `this_thread::sleep_for` command.